

Final Project Design Document

Overview:

Our client design is a simple loop that ingests user commands from the command line, packages them into a message format, sends them to the server it's connected to, then waits for a response from the server. To connect with a server, the client establishes a private two-way communication channel using spread. The details of how this is done are outlined in the next section.

Five servers are responsible for maintaining the state of this application. The servers will quickly reach consensus about the state of the application once they are connected to each other, although two servers that are partitioned may recognize different states. These states will be reconciled once the servers' information can reach each other.

This consensus algorithm in this design is based on an eventual path propagation algorithm. When there's a change in the membership of a group, all servers in the group exchange their view of the network, expressed through a summary 2D matrix of messages they believe have been received by the other servers. The server that has the most up to date messages from a given server is responsible for distributing those messages to the other servers.

The messages that change the state of the application are: new messages, read updates, and delete updates. For each server, an append only log is maintained to keep track of these updates. This is described in further detail below.

Group Membership:

Each server starts as a member of two groups. Server *i* will be a member of group `server_i_in` and group `all_servers`. `server_i_in` is used to establish a connection and send messages to server *i*, and only has server *i* as a member. `all_servers` is the group used for communication among servers.

Each client starts as a member of two groups, `client_<name>_connect` and `client_<name>_in` where `<name>` is a uniquely identifying integer generated at runtime.

Communicating between client and server:

When a client wants to connect to server *i*, it will send a `CONNECT` message to group `server_i_in` with the name of the connecting client. Upon reading this message, server *i* will join the group `client_<name>_connect`. The client will receive a membership message when this connection occurs, which is how it can confirm that a successful connection has been established. If server *i* crashes or the client leaves, the other party will get a membership message in `client_<name>_connect` notifying it of the absence.

As long as both parties are members of `client_<name>_connect`, the client is allowed to send messages to the server. To send a request, the client will send a message to group `server_i_in` with a command. The server will respond by sending a message to `client_<name>_in`. This is how a private two way communication channel can be created between server and client.

Logging

The primary method of persistence in our application is an append-only log for each server. For server `i` the file `server_i.log` on server `j` contains a list of all message, read, and delete updates known by server `j` to have originated from server `i`. The messages can be applied in any order to achieve the same state because they have a precedence of delete > read > message. That is, if a server tries to apply a read or delete update for a message it has not yet received, it must apply that update to the message upon receiving it. Similarly, if a read update and delete update are both applied to the same message, no matter the order the message will be deleted after both updates.

To serve user requests more quickly, the server stores the state in memory as a data structure which maps usernames to inboxes. This way the server does not need to apply all updates in a log in order to show a user her mailbox. We would also like to avoid the unscalable practice of applying all log updates to the state every time the server starts. To do this, we periodically serialize the local in-memory state to a file as a JSON object so that on startup, only the updates that were written after serializing the state need to be reapplied.

The final wrinkle in this logging procedure is that each file named `server_i.log` is actually a series of small files (named `server_i_X.log` where `X` is a monotonically increasing integer) so that the server can discard messages once they are no longer necessary. Once a server has knowledge that all other servers have received message `n` from a particular server, it can delete all log files for which the entries contained within the file have indices less than `n`.

Client protocol:

State:

Bool connected
Bool logged_in
string connected_server_group
String username

Int session id

list<MailMessage> inbox

Repeat:

 wait for input from user or message from spread

On receive user input:

 If command == 'c <new server>':

 If connected:

 Call Disconnect

 Set connected_server_group = server_<new server>_in

 Send ConnectMessage to connected_server_group with client id as data

 Set connection timeout

 If command == 'u <username>'

 If logged_in:

 Call Log out

 Set username = username

 If command == 'm'

 Require client is connected

 Require user is logged in

 Send MailMessage with user input as data to

 Connected_server_group

 Add message stamp as a key to requests

 If command == 'r <message_index>'

 Require client is connected

 Require user is logged in

 Require message index exists

 Send ReadMessage with inbox[message_index].id as data

 Wait to read server response

 If command == 'd <message_index>'

 Require client is connected

 Require user is logged in

 Require message index exists

 Send DeleteMessage with inbox[message_index].id as data

 Wait to read server response

 If command == 'l'

 Require client is connected

 Require user is logged in

 Send ShowInboxMessage

 Clear inbox

 Wait for server response

 If membership message is received from client_<name>_connect:

 Log out

 If ServerInboxResponse is received

 Append mail headers to inbox list

 Print mail headers

 If command == 'v'

 Require client is connected

- Require user is logged in
- Send ShowComponentMessage
- Wait for server response

While waiting for server response:

- On receive INBOX message:

 - Append contents to inbox

 - Print out contents

- On receive COMPONENT message:

 - Print out contents

- On receive membership message from client_<name>_connect:

 - Disconnect

 - Stop waiting

- On receive ACK message

 - Stop waiting

On receive membership message from client_<session_id>_connect

- If number of members == 2:

 - Set connected = true

- Else

 - Call Disconnect

 - Alert user that they must connect to a different server

On connection timeout:

- Call Disconnect

- Alert user that they must connect to a different server

Subroutine Disconnect:

- If connected:

 - Set connected = false

 - Leave group server_<session_id>_connect

 - Generate new session_id

 - Join group server_<session_id>_connect

Subroutine Log out:

- Set logged_in to false

Server protocol:

State:

int knowledge[5][5] where knowledge[i][j] is the last message the ith machine that it has received from machine j, determined by knowledge messages we have received.

```
UserCommand * log[5][variable]  
UserCommand * queue[MAX_QUEUE_LENGTH]  
Set<string> client_connections
```

State state

Algorithm:

Read file state.json (if exists) into boost::property_tree object, populate state with object's data

For each server n = 1 to 5:

```
    Read log file starting at line state.messages_safe_delivered[n]  
    While more lines to read:  
        If command index > state.applied_to_state[n]  
            Call apply_update(log update)  
        Append log update to log[this machine][n]
```

On receive CONNECT message in server_i_in:

```
    Add client name in message to client_connections  
    Join group client_<name>_connect
```

On receive membership message from client_<name>_connect

```
    If number of members < 2:  
        Leave group  
    Remove client_<name>_connect from client_connections
```

On receive M,R,D message in group server_i_in data:

```
    Increment knowledge[this machine][this machine]
```

```
    Extract message info into UserCommand command  
    Write command to log file  
    Call apply_update(command)
```

```
    Increment state.applied_to_state[this machine]  
    Send command to group
```

On receive user command in group all_servers:

```
    If message index == knowledge[this machine][message origin] + 1  
        Write command to log file  
        Increment knowledge[this machine][message origin]  
        Call apply_update(command)  
        Increment state.applied_to_state[message origin]
```

```
    If state.applied_to_state % knowledge_timeout == 0
        Send knowledge message to all_servers
```

On receive knowledge message from server i in group all_servers:

```
    Replace row i in knowledge with message content
```

```
    Call collect_garbage()
```

On receive membership message from all_servers

```
    Go to state synchronize
```

Synchronize:

```
    Send knowledge message
```

```
    Wait for knowledge message from all others in group
```

```
        If receive M,R,D message in group server_i_in data:
```

```
            Write message to log file
```

```
            Push message into queue
```

```
    Set S to be server indices in this group
```

```
    For i = 0 to 4
```

```
        If this server has the largest value in column i of knowledge
```

```
            Send messages starting from min(column i of knowledge among servers
            in S) to max(column i of knowledge among servers in S)
```

```
    If message are safe to discard, discard and increment state.message_safe_delivered
```

```
    Pop queued messages one by one and process them as if a user just sent them.
```

Subroutine collect_garbage(i):

```
    State.safe_delivered[i] = min element in column i of knowledge
```

```
    While first entry in log[i] is valid message and minimum entry in column i of knowledge >
    log[i][0]:
```

```
        Pop front of log[i]
```

```
    Delete all log files from machine i that are full and only contain commands with index \
    less than state.safe_delivered[i]
```

Subroutine serialize_state

```
    Convert state to boost::property_tree object
```

```
    Write object to file state.json
```

Subroutine apply_update(UserCommand command)

```
    If type == DELETE
```

```
        If message in inboxes[uid]:
```

```
            Remove message from inboxes[uid]
```

```
            append command.id to state.deleted
```

```
        Else:
```

```
            If command.id in state.pending_read:
```

```
                Remove command.id from state.pending_read
```

```

        Append command.id to state.pending_delete
    If type == READ
        If message in inboxes[uid]:
            Mark message in inboxes[uid] as read
        Else if command.id in state.pending_delete or state.deleted:
            Do nothing
        Else:
            Append command.id to state.pending_read

    If type == MAIL:
        If command.id in state.pending_delete:
            remove command.id from state.pending_delete
            Put command.id in state.deleted
            return
        Else if command.id in state.pending_read:
            Remove command.id from state.pending_read
            Extract command info into MailMessage m
            Mark m as read
        Else:
            Extract command info into MailMessage m
            Insert m into inboxes[uid]
    Periodically serialize state and broadcast knowledge message to network

```

Data Structures

```

enum MessageType
{
    // Client to server messages
    CONNECT,
    MAIL,
    READ,
    DELETE,
    SHOW_INBOX,
    SHOW_COMPONENT,

    // Server to client message
    ACK,
    INBOX,
    RESPONSE,
    COMPONENT,

    // Server to server messages

```

```

        COMMAND,
        KNOWLEDGE
};

/* Unique identifier for any command */
struct MessageIdentifier
{
    int index;
    int origin;
};

/* Message signaling incoming client connection */
struct ConnectMessage
{
    MessageType type = MessageType::CONNECT;
    uint32_t session_id;
};

/* Message containing new mail message from client */
struct MailMessage
{
    MessageType type = MessageType::MAIL;
    uint32_t session_id;
    int seq_num;
    char username[MAX_USERNAME];
    char to[MAX_USERNAME];
    char subject[MAX_SUBJECT];
    char message[EMAIL_LEN];
};

/* Message requesting mail to be marked as read by client */
struct ReadMessage
{
    MessageType type = MessageType::READ;
    uint32_t session_id;
    int seq_num;
    char username[MAX_USERNAME];
    MessageIdentifier id;
};

```



```

/* Message requesting mail be deleted by client */
struct DeleteMessage
{
    MessageType type = MessageType::DELETE;
    uint32_t session_id;
    int seq_num;
    char username[MAX_USERNAME];
    MessageIdentifier id;
};

/* Message sent by client to request the current user's inbox */
struct GetInboxMessage
{
    MessageType type = MessageType::SHOW_INBOX;
    uint32_t session_id;
    int seq_num;
    char username[MAX_USERNAME];
};

/* Message sent by client to request all connected server names */
struct GetComponentMessage
{
    MessageType type = MessageType::SHOW_COMPONENT;
    uint32_t session_id;
};

/* Wrapper for Read, Mail, and Delete messages with metadata used by
Server */
struct UserCommand
{
    MessageIdentifier id;
    time_t timestamp;
    std::variant<
        MailMessage,
        ReadMessage,
        DeleteMessage
    > data;
};

/* Message used for sending a server's network view to other servers */

```

```

struct KnowledgeMessage
{
    MessageType type = MessageType::KNOWLEDGE;
    int sender;
    int summary[N_MACHINES][N_MACHINES];
};

/* Message to end response to a client */
struct AckMessage
{
    MessageType type = MessageType::ACK;
    int seq_num;
    char body[300];
};

/* Struct containing all critical user data about a mail message */
struct InboxEntry
{
    bool read;
    time_t date_sent;
    char to[MAX_USERNAME];
    char from[MAX_USERNAME];
    char subject[MAX_SUBJECT];
    char message[EMAIL_LEN];
};

/* Struct that stores user and server data for a mail message */
struct InboxMessage
{
    MessageType type = MessageType::INBOX;
    MessageIdentifier id;
    InboxEntry msg;
};

/* Message used to share the names of current connected group with client
*/
struct ComponentMessage
{
    int num_servers;
    char names [5][MAX_GROUP_NAME];
};

```

```

};

/* Wrapper for any server message to client */
struct ServerResponse
{
    MessageType type = MessageType::RESPONSE;
    int seq_num;
    std::variant<
        AckMessage,
        InboxMessage,
        ComponentMessage
    > data;
};

/* Struct containing data displayed by client after an inbox request */
struct InboxHeader
{
    time_t timestamp;
    char subject [MAX_SUBJECT];
    char sender [MAX_USERNAME];
    bool read;
    MessageIdentifier id;
};

/* Message containing up to 20 inbox headers */
struct ServerInboxResponse
{
    MessageType type = MessageType::RESPONSE;
    int mail_count;
    InboxHeader inbox[20];
};

/*
    In memory state. This is what gets serialized to disk periodically.

```

```

*/
struct State
{
    int knowledge[N_MACHINES][N_MACHINES];
    int safe_delivered[N_MACHINES];
    int applied_to_state[N_MACHINES];
    std::unordered_map<std::string, std::multiset<InboxMessage>> inboxes;
    std::set<MessageIdentifier> pending_delete;
    std::set<MessageIdentifier> pending_read;
};

```

Scenarios that have been tested:

If no additional note, means it behaved as expected during testing.

List, mail, read, delete, component with multiple servers connected, changes synced

RECOVERY/CRASHING

State recovery with multiple servers (1-5 synced, mail from s1, kill all, then only bring back 5, mail still there)

Start all servers, crash one, action, recover

PARTITIONING

partition, read, merge

Partition, mail, merge

Partition, read & delete same message, merge

Partition, mail + read + delete, partition, mail + read + delete, merge (does not work, read and inbox is all wrong)

*Occasional bug with listing components, number is right but server not printed after partitioning
- STASHED

*Cant connect to a partitioned server - FIXED

*mail not displayed in time sorted order - FIXED

*Occasionally connecting with another client will disconnect some other clients/itself

Garbage collection - IN PROGRESS

Periodically send knowledge messages - IN PROGRESS

Crash before completely sending synch

- No mail being sent
- Pending delete may not work